

Aman Saxena

8840154168 | aman16nov.as@gmail.com | GitHub | LinkedIn | LeetCode | Portfolio

TECHNICAL SKILLS

Languages: C/C++, Python, JavaScript, TypeScript

Frontend: React, Redux Toolkit (RTK), HTML, CSS, Tailwind CSS

Backend: Django, FastAPI, Flask, REST APIs, Celery, Kafka, Redis, WebSocket

Databases: PostgreSQL, SQL, Materialized Views, Query Optimization & Indexing

System Design: Microservices, Event-Driven Architecture, Caching Strategies, Message Queues, Load Balancing, Database Sharding & Replication

Cloud: AWS

Gen AI : LangChain, LLM APIs, RAG, Vector Databases, AI Agents

Tools & Platforms: Git, GitHub, Docker, Postman, Linux, Celery, Kafka, Redis, WebSocket

Deployment & DevOps: Docker, Gunicorn, Render, Vercel, CI/CD , Container Memory & Cold-Start Optimisation

Concepts: OOPs, Data Structures & Algorithms, RBAC / Access Control

EXPERIENCE

Einstellen.AI

Noida, India

Full Stack Developer (Product-Based)

Feb 2025 – July 2026

- Built an AI-driven hiring platform using Django and a microservices architecture, owning backend design and development across 4 services.
- Designed a feature-level access-control (RBAC) system across the Django microservices and React/TypeScript frontend, using Redis-cached, fail-closed authorization checks for consistent, low-latency permission handling.
- Built asynchronous AI-based interview workflows and automated report generation using Celery and REST APIs, improving system responsiveness under concurrent load.
- Built reusable UI components in TypeScript and React, and used Redux Toolkit (RTK) for state management across the frontend.
- Added real-time chat using WebSockets for bi-directional communication between users on the platform.
- Used Apache Kafka to decouple microservices through an event-driven architecture, improving scalability and fault isolation.
- Improved PostgreSQL performance for large, structured datasets through query tuning, indexing, and materialized views, reducing data retrieval latency.
- Built a Resume AI Parser using the Ollama LLM API to extract structured fields (skills, experience, education) from candidate resumes and generate a match score against job descriptions, automating an initial layer of candidate screening.

Codezin Technology Solutions

Kolkata, India

Full Stack Developer Intern

Aug 2023 – Oct 2023

- Developed dynamic, scalable web applications using Django, HTML, and JavaScript.
- Revamped interactive website features that increased user engagement by 30%.
- Led cross-browser compatibility testing, achieving 95% usability consistency across devices.

PROJECTS

Multi-Channel Notification Fan-Out System | Django, Kafka, Celery, Redis, PostgreSQL

- Built an event-driven notification system that fans out a single business event (e.g., "order shipped") to Email, SMS, and Push channels via Kafka topics and independent Celery workers, so one channel's failure doesn't block the others.
- Used Redis distributed locks and database-level unique constraints to keep message processing idempotent, avoiding duplicate sends under Kafka's at-least-once delivery guarantees.
- Added a retry mechanism with exponential backoff and a dead-letter queue for failed deliveries, so provider outages don't affect the rest of the pipeline.
- Built a PostgreSQL schema with a materialized view tracking per-channel delivery success rates for real-time visibility into notification health.

MyTwinAI | Django, DRF, React, TypeScript, Chroma, ONNX Runtime, Groq, Docker | Live Link

- Built an AI-powered personal portfolio assistant (solo project) that answers recruiter questions about my experience, skills, and projects by retrieving from my own structured profile data via a RAG pipeline, reusing and adapting the architecture from my RAG_Chat_App project.
- Streamed LLM responses token-by-token over Server-Sent Events and auto-generated 3 relevant follow-up questions after each reply to guide conversation.
- Cached the most common Q&A pairs to serve them in milliseconds at zero LLM token cost, and added automatic fallback to a secondary LLM when the primary model's daily token limit is hit, so the app degrades gracefully instead of failing.
- Dockerized the Django/DRF backend and React/TypeScript frontend and deployed them independently on Render and Vercel.

EDUCATION

GL Bajaj Institute of Technology, AKTU University

Greater Noida, India

Master of Computer Applications (MCA)

2023 – 2025